

# Dimensionality Reduction using Principal Component Analysis

Sina Tootoonian

# Outline

- 1 Motivation
- 2 Variance of single neurons
- 3 Variance of pseudo-neurons
- 4 Some Facts about Matrices
- 5 Singular Value Decomposition
- 6 Maximum Variance Directions
- 7 Adjacent Approaches [3/3]
- 8 **TODO** Summary

# Section 1

## Motivation

# Neuroscience data arrive high dimensional

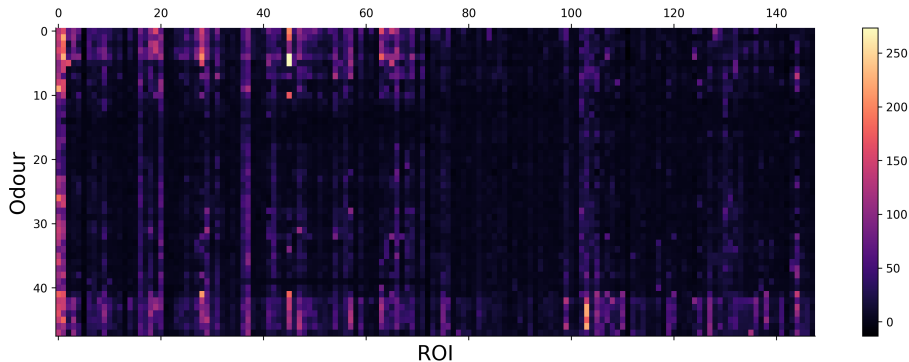


Figure: Response of 150 olfactory glomeruli to 48 odours (Tobias Ackels)

# Visualization requires 2-3 dimensions

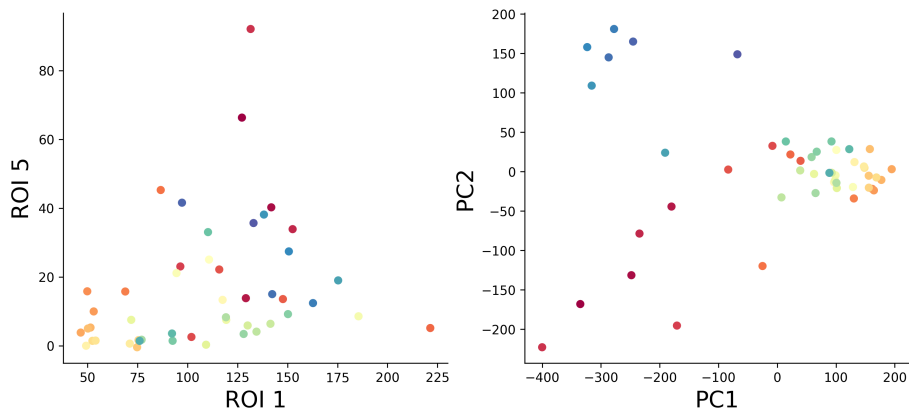


Figure: Odour responses of two ROIs, vs. two principal components.

# High-D data can be embeddings of low-D manifolds

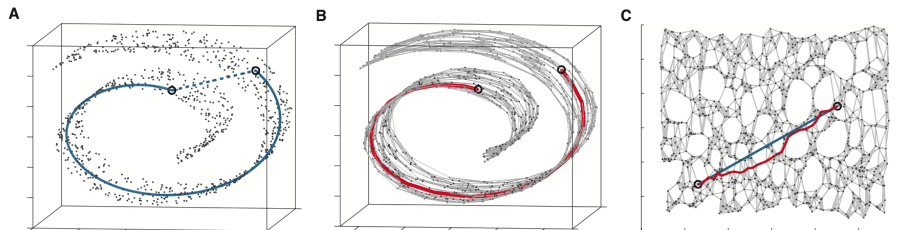
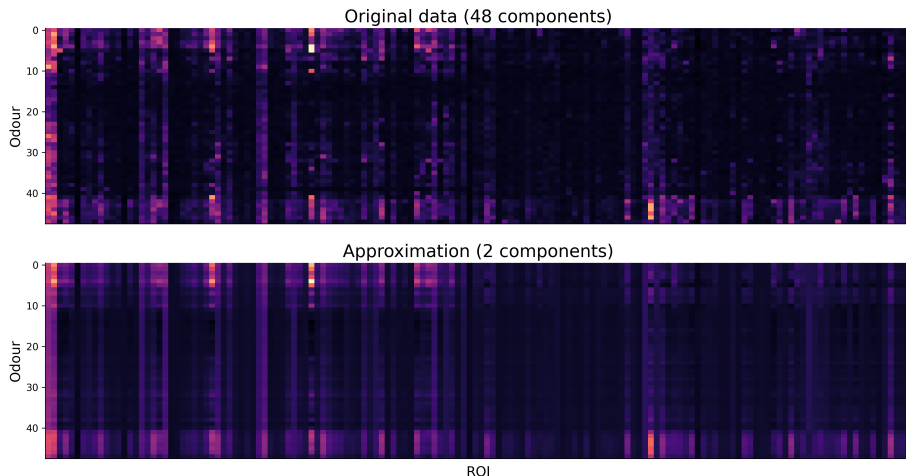


Figure: 'Swiss roll' dataset from Tenenbaum et al. 2000.

# Data compression / approximation / denoising



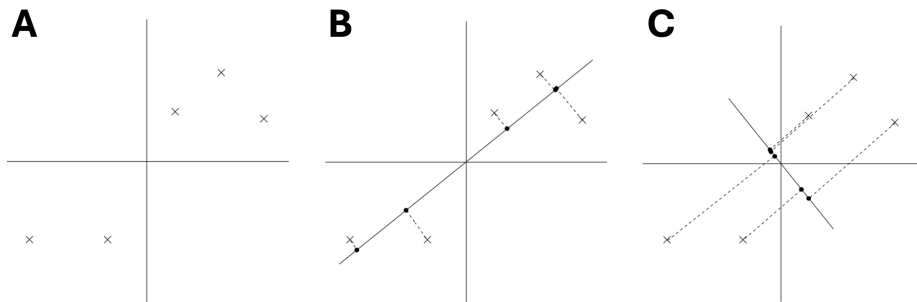
**Figure:** Raw odour responses and 2-component approximation.

Find **directions** in data space that **maximize variance**.

- What do we mean by *directions*?
- What do we mean by *variance*?
- Why is *maximizing* a good idea?



# Key idea behind PCA



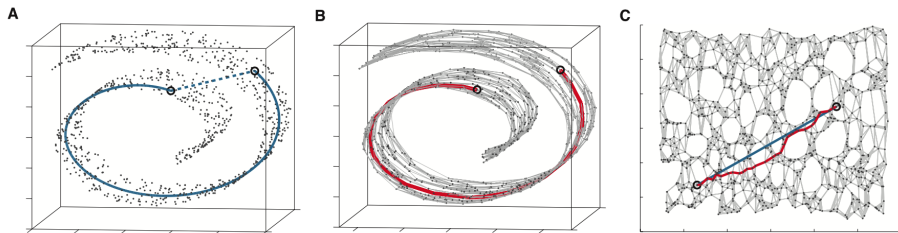
**Figure:** Projecting data (A) along directions that capture much (B) or little (C) variance. From CS 229.

- Coordinate systems for data
- Neurons and Pseudo-neurons
- Measuring activity of single neurons with variance
- Measuring co-activity of neurons with covariance
- Relating variance of pseudo-neurons to covariance

## Section 2

### Variance of single neurons

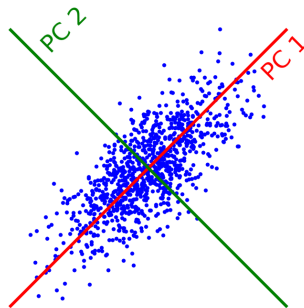
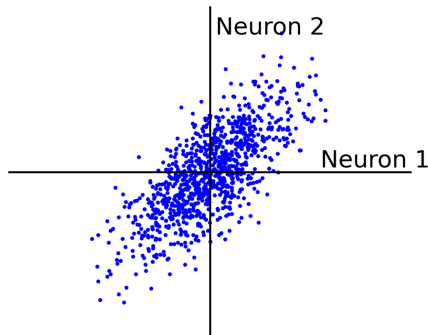
# Notions of dimensionality



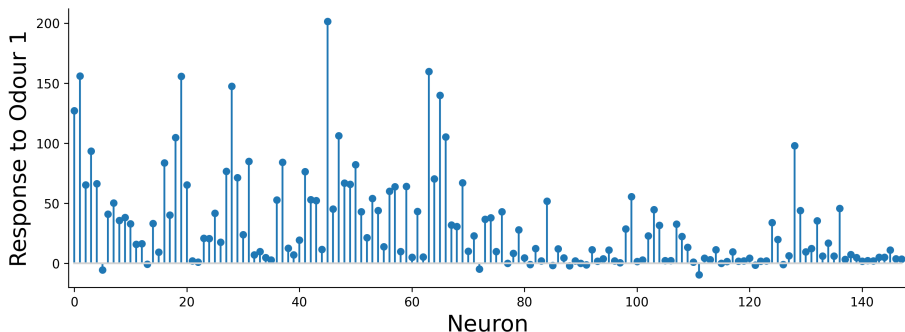
- Extrinsic dimension
- Intrinsic dimension
  - Lower because of correlations in the data
- PCA: Find the intrinsic, **linear** dimension

# The importance of coordinate systems

- Laws of physics don't depend on coordinates.
- Laws of neuroscience depends on having the right coordiantes
  - Usually not the ones your data is recorded in.
- PCA finds coordinates that are **matched to the data**.



# Standard coordinates

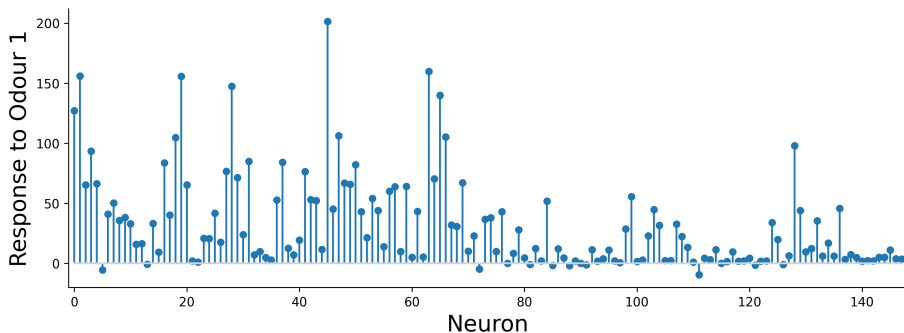


- Notice implicit coordinates:

$$\begin{aligned}\mathbf{x} &= [x_1, x_2, \dots] \\ &= x_1 \mathbf{e}_1 + x_2 \mathbf{e}_2 + \dots\end{aligned}$$

- Each coordinate  $\mathbf{e}_1, \mathbf{e}_2$  corresponds to unit activity of one neuron.

# Coordinates are orthonormal



$$\mathbf{x} = x_1 \mathbf{e}_1 + x_2 \mathbf{e}_2 \dots$$

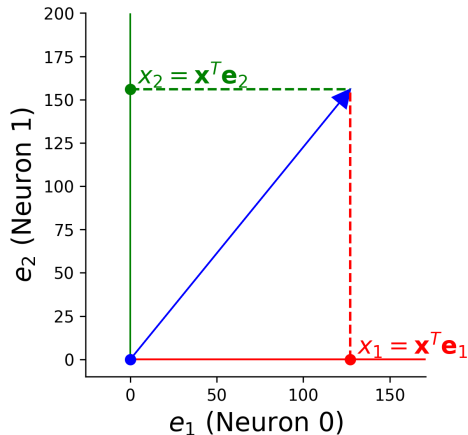
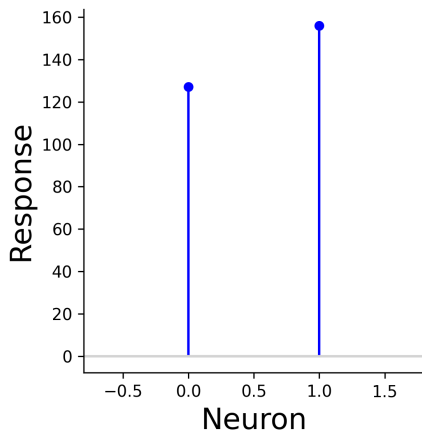
- The coordinate system  $\{\mathbf{e}_1, \mathbf{e}_2, \dots\}$  is:
  - Complete: We can represent any activity vector in it.
  - Ortho...

$$\mathbf{e}_1^T \mathbf{e}_2 = 0, \dots$$

- ...normal

$$\mathbf{e}_1^T \mathbf{e}_1 = 1, \quad \mathbf{e}_2^T \mathbf{e}_2 = 1, \dots$$

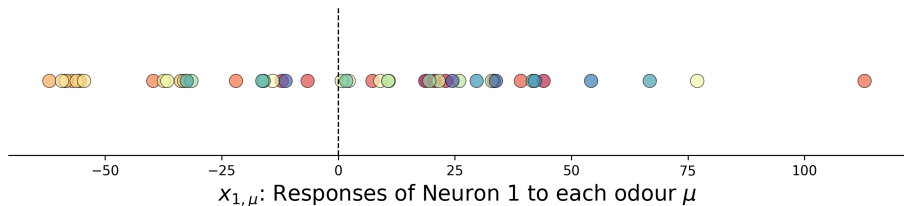
# Extracting single unit responses by projection



$$x_1 = \mathbf{x}^T \mathbf{e}_1 = \mathbf{e}_1^T \mathbf{x}$$



# Different ways of summarizing activity



- Mean?

$$\bar{x}_1 = \frac{1}{\# \text{ stimuli}} \sum_{\mu} x_{1,\mu} = \langle x_{1,\mu} \rangle.$$

- Absolute value?

$$\overline{|x_1|} = \langle |x_{1,\mu}| \rangle$$

- Absolute value relative to mean?

$$\overline{|x_1 - \bar{x}_1|} = \langle |x_{1,\mu} - \bar{x}_1| \rangle.$$

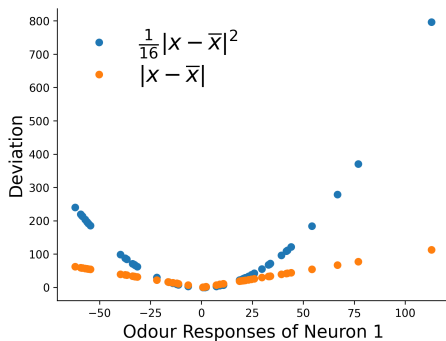
- Squared value?

$$\overline{x_1^2} = \langle x_{1,\mu}^2 \rangle$$

# Variance of a single neuron

$$\text{var}(x_1) = \langle (x_{1,\mu} - \bar{x}_1)^2 \rangle.$$

- Average energy relative to the mean
- Mathematically tractable ✓
- Susceptible to outliers ✗



## Section 3

### Variance of pseudo-neurons

# Other orthonormal coordinates

We can describe the same activity in terms of ‘pseudo-neurons.’

$$\mathbf{x} = \tilde{x}_1 \tilde{\mathbf{e}}_1 + \tilde{x}_2 \tilde{\mathbf{e}}_2 + \dots$$

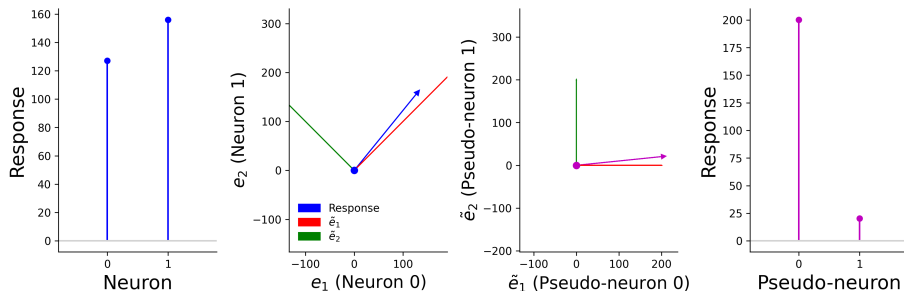
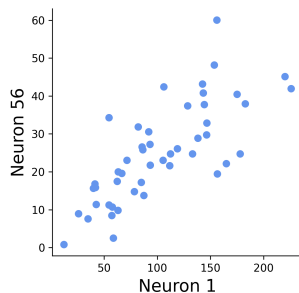
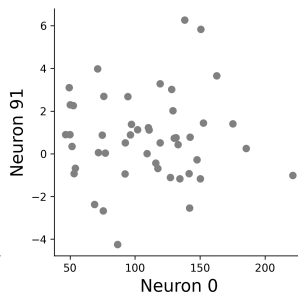
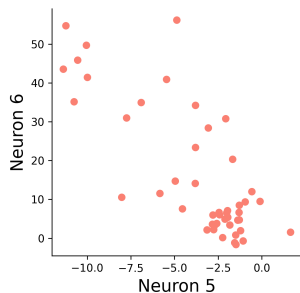


Figure: Responses of neurons and pseudo-neurons to the first odour.

# Covariance of neural populations

- Neurons don't respond independently, but frequently **covary**



- **C**ovariance measures covariation of a neuron with another:

$$\text{cov}(x_1, x_2) = \langle (x_{1,\mu} - \bar{x}_1)(x_{2,\mu} - \bar{x}_2) \rangle.$$

- Variance is covariation of a neuron with itself!

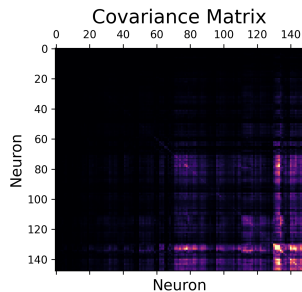
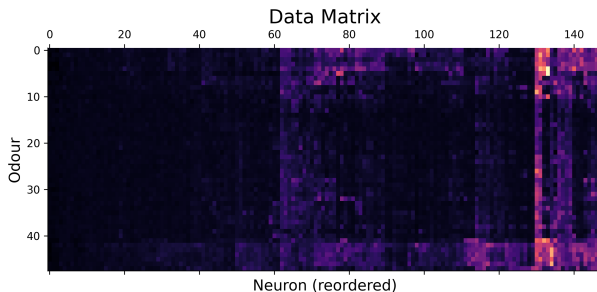
$$\begin{aligned}\text{var}(x_1) &= \langle (x_{1,\mu} - \bar{x}_1)^2 \rangle \\ &= \langle (x_{1,\mu} - \bar{x}_1)(x_{1,\mu} - \bar{x}_1) \rangle.\end{aligned}$$

# Covariance matrix

- The **covariance matrix** tabulates covariance for all pairs of neurons.

$$\text{cov}(\mathbf{x}) = \begin{bmatrix} \text{var}(x_1) & \text{cov}(x_1, x_2) & \dots \\ \text{cov}(x_1, x_2) & \text{var}(x_2) & \dots \\ \vdots & \vdots & \ddots \end{bmatrix} = \langle (\mathbf{x}_\mu - \bar{\mathbf{x}})(\mathbf{x}_\mu - \bar{\mathbf{x}})^T \rangle$$

- Diagonals have variances
- Off-diagonals have covariances



- Not just useful book keeping. . .

# Variance of a single neuron (again)

- Previously we just 'took' the data for neuron 1.
- We can view this as first projecting along the first coordinate:

$$x_{1,\mu} = \mathbf{x}_\mu^T \mathbf{e}_1$$

- Then computing variance

$$\text{var}(x_1) = \langle (x_{\mu,1} - \bar{x}_1)^2 \rangle$$

# Variance of one pseudoneuron

- Activity of a **pseudoneuron**:

$$\tilde{x}_{1,\mu} = \mathbf{x}_\mu^T \mathbf{u}_1$$

- Mean activity of a pseudoneuron:

$$\begin{aligned}\bar{\tilde{x}}_1 &= \langle \mathbf{x}_\mu^T \mathbf{u}_1 \rangle \\ &= \langle \mathbf{x}_\mu^T \rangle \mathbf{u}_1 \\ &= \bar{\mathbf{x}}^T \mathbf{u}_1.\end{aligned}$$

- Variance of a pseudoneuron

$$\begin{aligned}\text{var}(\tilde{x}_1) &= \langle (\tilde{x}_1 - \bar{\tilde{x}}_1)^2 \rangle \\ &= \langle (\mathbf{x}_\mu^T \mathbf{u}_1 - \bar{\mathbf{x}}^T \mathbf{u}_1)^2 \rangle\end{aligned}$$



# Variance of *any* pseudoneuron

- We don't have to go through the procedure again!
- Variance of any pseudoneuron  $\tilde{x} = \mathbf{x}^T \mathbf{u}$  is

$$\text{var}(\tilde{x}) = \mathbf{u}^T \text{cov}(\mathbf{x}) \mathbf{u}.$$

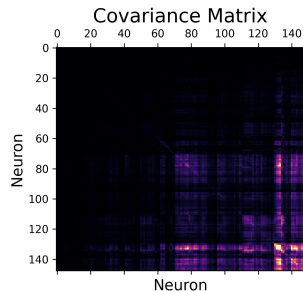
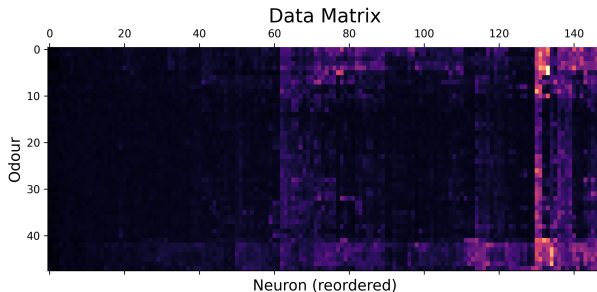
- PCA now becomes finding the  $\mathbf{u}$  that maximizes this variance.
- How do we do this?

## Section 4

### Some Facts about Matrices

# Matrices

- Some matrices we've already encountered:



- Data matrix (rectangular)
- Covariance matrix (square, symmetric)
  - Why is it symmetric?

# Different ways to view matrices

$$\mathbf{A} = \underbrace{\begin{bmatrix} A_{11} & A_{12} & \dots \\ A_{21} & A_{22} & \dots \\ \vdots & \vdots & \ddots \end{bmatrix}}_{\text{Table of elements}} = \underbrace{\begin{bmatrix} \mathbf{r}_1^T \\ \mathbf{r}_2^T \\ \vdots \\ \mathbf{r}_M^T \end{bmatrix}}_{\text{Stacked rows}} = \underbrace{\begin{bmatrix} \mathbf{c}_1 & \mathbf{c}_2 & \mathbf{c}_3 & \dots & \mathbf{c}_N \end{bmatrix}}_{\text{Stacked columns}}.$$

# Matrix operations

- **Linearly** transform N-dimensional inputs  $\mathbf{x}$  into M-dimensional outputs  $\mathbf{y}$ ,

$$\mathbf{y} = \mathbf{A}\mathbf{x}.$$

- Can think of this element-wise:

$$y_i = \sum_{j=1}^N A_{ij} x_j.$$

- Can think of this as projecting  $\mathbf{x}$  on each row,

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_M \end{bmatrix} = \begin{bmatrix} \mathbf{r}_1^T \mathbf{x} \\ \mathbf{r}_2^T \mathbf{x} \\ \vdots \\ \mathbf{r}_M^T \mathbf{x} \end{bmatrix}.$$

- Can think of this as summing the columns, weighted by  $\mathbf{x}$ ,

$$\mathbf{y} = \sum_{i=1}^N \mathbf{c}_i x_i.$$

# Example Matrices

Name	Matrix $A$	Action $y = Ax$
Zero	$\begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix}$	$y = 0$
Identity	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$	$y = x$
All ones	$\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$	$y = \begin{bmatrix} \sum_i x_i \\ \sum_i x_i \end{bmatrix}$
Uniform scaling	$\begin{bmatrix} k & 0 \\ 0 & k \end{bmatrix}$	$y = \begin{bmatrix} kx_1 \\ kx_2 \end{bmatrix}$
Diagonal	$\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$	$y = \begin{bmatrix} ax_1 \\ bx_2 \end{bmatrix}$
Permutation	$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$	$y = \begin{bmatrix} x_2 \\ x_1 \end{bmatrix}$
Rotation	$\begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$	$y$ is $x$ rotated by $\theta$ .

# Composing transformations

- We can form complex transformations by composing simple ones.
- For example, a scaling and a rotation:

$$\mathbf{y} = \underbrace{\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}}_{\text{scaling}} \underbrace{\begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}}_{\text{rotation}} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \underbrace{\mathbf{D}\mathbf{R}}_{\mathbf{A}} \mathbf{x}.$$

## Section 5

# Singular Value Decomposition



# All matrices are diagonal matrices (in the right coordinates)

- Diagonal matrices were easy to work with

$$\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} ax_1 \\ bx_2 \end{bmatrix}$$

- What about an arbitrary matrix? Looks complex. . .

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} \sum_j A_{1j}x_j \\ \sum_j A_{2j}x_j \end{bmatrix}.$$

- Surprise: Every matrix  $\mathbf{A}$  is the composition of just three operations!

$$\mathbf{A} = \underbrace{\mathbf{U}}_{\text{rotate}} \underbrace{\mathbf{S}}_{\text{scale}} \underbrace{\mathbf{V}^T}_{\text{project}}.$$

# Three parts of Singular Value Decomposition

$$\mathbf{A} = \underbrace{\mathbf{U}}_{\text{rotate}} \underbrace{\mathbf{S}}_{\text{scale}} \underbrace{\mathbf{V}^T}_{\text{project}} .$$

- Columns of  $\mathbf{V}$  form orthonormal coordinates for the **input** space.
- Columns of  $\mathbf{U}$  form orthonormal coordinates for the **output** space
- Diagonal matrix  $\mathbf{S}$  of non-negative **singular values** apply a scaling.
- If we:
  - Use  $\mathbf{V}$  coordinates for the input, and
  - Use  $\mathbf{U}$  coordinates for the output, then
  - $\mathbf{A}$  is a scaling!

# Three parts of Singular Value Decomposition

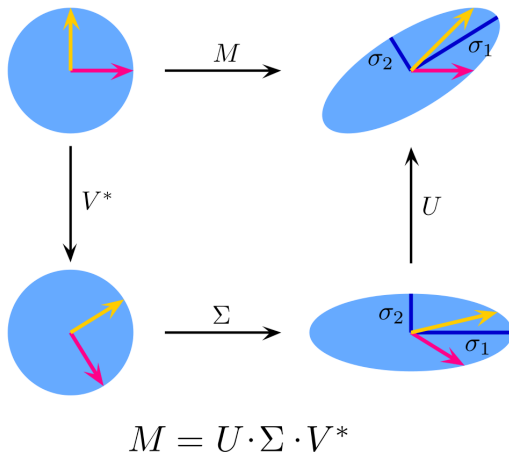


Figure: Three parts of Singular Value Decomposition (Wikipedia).

# Three steps of Singular Value Decomposition

$$\mathbf{Ax} = \underbrace{\mathbf{U}}_{\text{rotate}} \underbrace{\mathbf{S}}_{\text{scale}} \underbrace{\mathbf{V}^T}_{\text{project}} \mathbf{x}.$$

- ❶ Project  $\mathbf{x}$  onto the input coordinates:

$$\mathbf{V}^T \mathbf{x} = \begin{bmatrix} \mathbf{v}_1^T \mathbf{x} \\ \vdots \\ \mathbf{v}_N^T \mathbf{x} \end{bmatrix} = \begin{bmatrix} \tilde{x}_1 \\ \vdots \\ \tilde{x}_N \end{bmatrix}$$

- ❷ Scale by  $\mathbf{S}$ :

$$\mathbf{SV}^T \mathbf{x} = \begin{bmatrix} s_1 & 0 & \dots \\ 0 & s_2 & \dots \\ \vdots & \vdots & \ddots \end{bmatrix} \begin{bmatrix} \tilde{x}_1 \\ \vdots \\ \tilde{x}_N \end{bmatrix} = \begin{bmatrix} s_1 \tilde{x}_1 \\ s_2 \tilde{x}_2 \\ \vdots \\ s_N \tilde{x}_N \end{bmatrix}$$

- ❸ Project out using the output coordinates

$$\mathbf{USV}^T \mathbf{x} = [\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_N] \begin{bmatrix} s_1 \tilde{x}_1 \\ s_2 \tilde{x}_2 \\ \vdots \\ s_N \tilde{x}_N \end{bmatrix} = \mathbf{u}_1 s_1 \tilde{x}_1 + \mathbf{u}_2 s_2 \tilde{x}_2 + \dots$$

# SVD of simple matrices

Name	A	U	s	V
Zero	$\begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix}$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$	$[0, \ 0]$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Identity	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$	$[1, \ 1]$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Negation	$\begin{bmatrix} -1 & 0 \\ 0 & -1 \end{bmatrix}$	$\begin{bmatrix} -1 & 0 \\ 0 & -1 \end{bmatrix}$	$[1, \ 1]$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
All ones	$\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$	$\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix}$	$[2, \ 0]$	$\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ -1 & 1 \end{bmatrix}$
Diagonal	$\begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$	$[2, \ 3]$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Permutation	$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$	$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$	$[1, \ 1]$	$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$
Rotation by $\theta$	$\begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$	$\begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$	$[1, \ 1]$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$

# SVD of covariance matrices

- Remember why we care: we're after the variance of pseudoneurons

$$\mathbf{u}^T \text{cov}(\mathbf{x}) \mathbf{u}.$$

- For covariance matrices, the input and output coordinates are the same

$$\text{cov}(\mathbf{x}) = \mathbf{V} \mathbf{S} \mathbf{V}^T$$

- Equalizer: Inputs are analyzed in  $\mathbf{V}$  coordinates and scaled.

$$\text{cov}(\mathbf{x}) \mathbf{u} = \sum_i \mathbf{v}_i \underbrace{s_i}_{\text{scale}} \underbrace{\mathbf{v}_i^T \mathbf{u}}_{\text{project}}$$

## Section 6

# Maximum Variance Directions

# Maximum variance direction from SVD

- We can use SVD to read-off the maximum variance direction(s) we need!
- Variance along a direction  $\mathbf{u}$

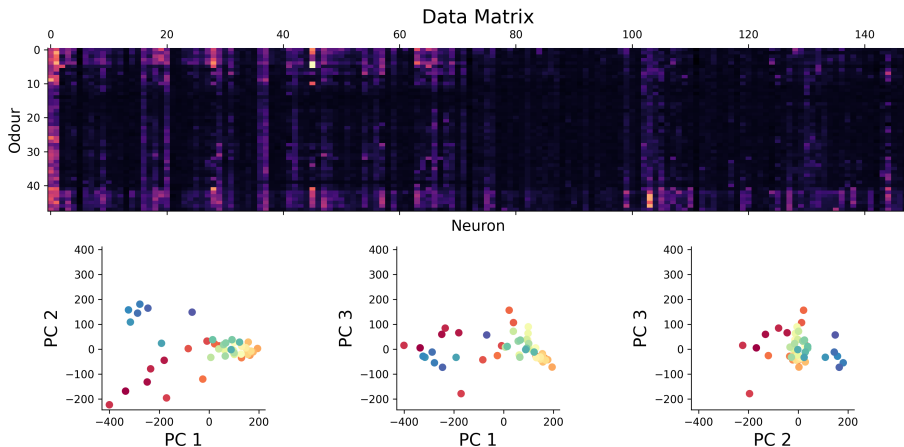
$$\begin{aligned}\mathbf{u}^T \text{cov}(\mathbf{x}) \mathbf{u} &= \mathbf{u}^T \underbrace{\left( \sum_i \mathbf{v}_i s_i \mathbf{v}_i^T \right)}_{\text{SVD}} \mathbf{u} \\ &= \sum_i (\mathbf{u}^T \mathbf{v}_i) s_i (\mathbf{v}_i^T \mathbf{u}) \\ &= \sum_i s_i (\mathbf{v}_i^T \mathbf{u})^2\end{aligned}$$

- Maximum variance direction is  $\mathbf{v}_1$
- Next highest variance direction is  $\mathbf{v}_2$ , etc.



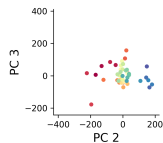
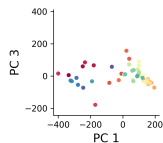
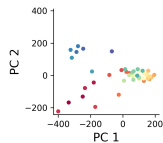
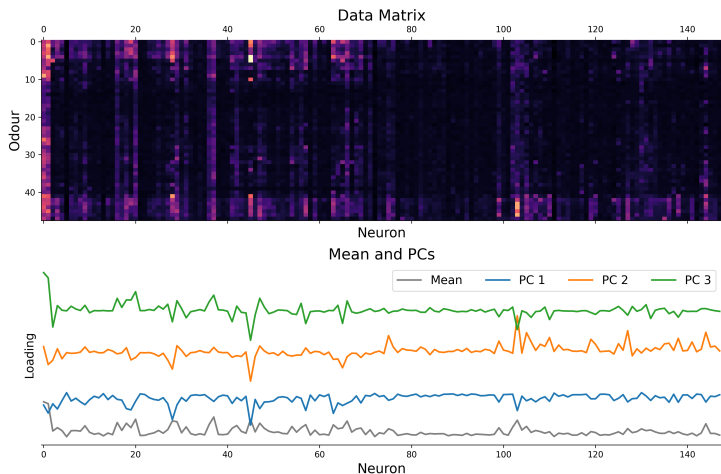
# Finally: Dimension Reduction with PCA

- Project data onto maximum variance directions



- Notice: projections are decorrelated

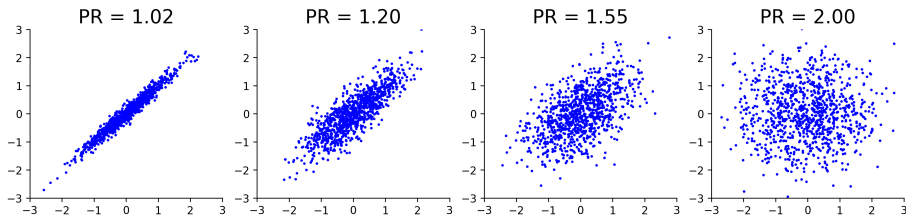
# Examining the Principal Components



# Measuring dimensionality with Participation Ratio

- Variances tell us energy in each direction
- Use this as a measure of dimensionality

$$PR = \frac{(\sum_i s_i)^2}{\sum_i s_i^2}.$$



# Approximation/Denoising with PCA

- Approximate using first  $K$  projections

$$\mathbf{x} \approx \underbrace{\sum_{i=1}^K (\mathbf{x}^T \mathbf{v}_i) \mathbf{v}_i}_{\text{Exact}} + \underbrace{\sum_{i=K+1}^D (\bar{\mathbf{x}}^T \mathbf{v}_i) \mathbf{v}_i}_{\text{Approximation}}.$$
$$\approx \bar{\mathbf{x}} + \sum_{i=1}^K (\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{v}_i \mathbf{v}_i$$

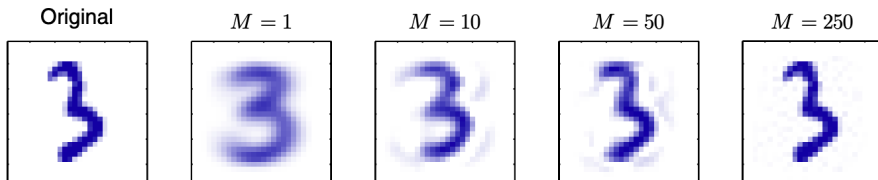


Figure: Approximating digits data using PCA (Bishop Fig 12.5)

# Approximation/Denoising with PCA

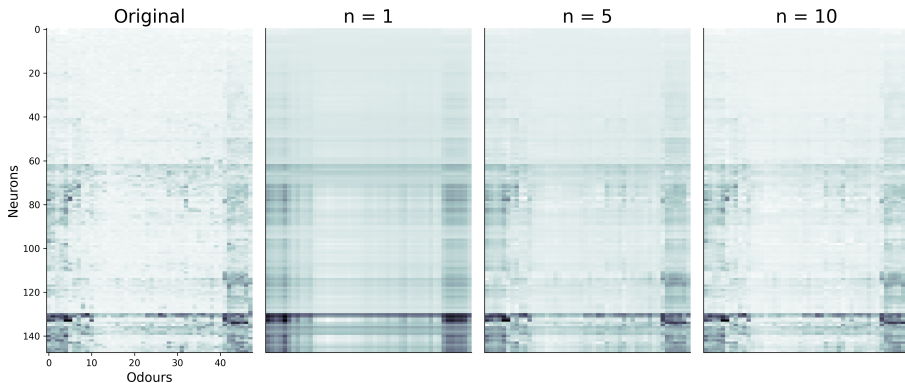


Figure: Approximating odour responses using PCA.

# How many dimensions to keep?

- The singular values tell us how much variance is explained by each dimension.
- We can use this to decide how many dimensions to keep.
- **Explained variance** measures the fraction of variance explained by the first  $K$  dimensions:

$$\text{EV}(K) = \frac{\sum_{i=1}^K s_i}{\sum_{i=1}^D s_i}.$$

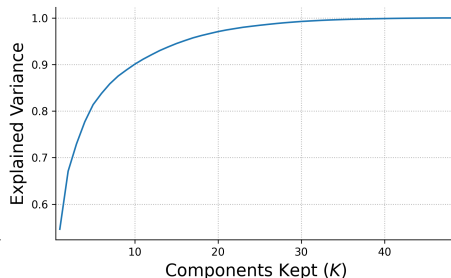
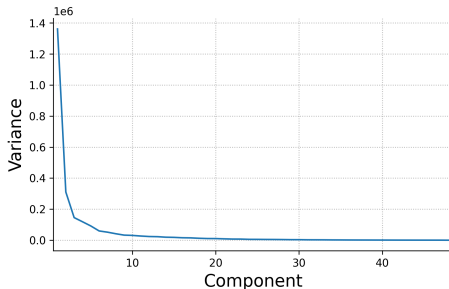


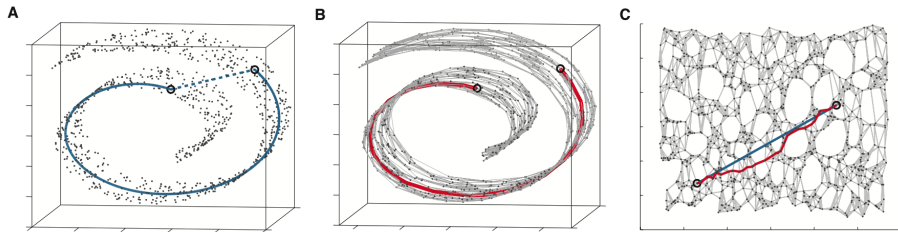
Figure: Explained variance for odour responses dataset.

## Section 7

### Adjacent Approaches [3/3]

# Exploiting nonlinearity with Kernel PCA

- PCA can be expressed in terms of similarity  $k(\mathbf{x}, \mathbf{y})$  between data points.
- PCA uses linear similarity  $k(\mathbf{x}, \mathbf{y}) = \mathbf{x}^T \mathbf{y}$ .
- Kernel PCA generalises this to allow other similarity measures.
- Nonlinear measures are sometimes appropriate.

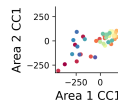
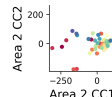
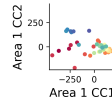
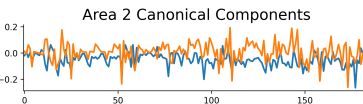
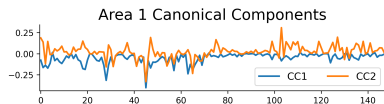
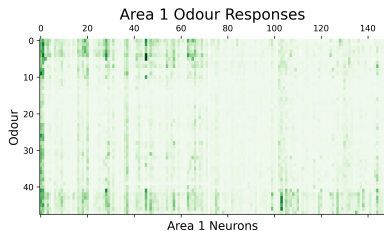


**Figure:** ISOMAP computes similarity as distance on the manifold.



# Comparing different datasets with CCA

- PCA finds maximum variance directions in one dataset
- CCA finds maximum co-variance directions in two datasets



# Supervised learning with LDA

- PCA doesn't care about class labels.
- Maximum variance isn't always best for discrimination.
- LDA: Finds directions that best discriminate data.

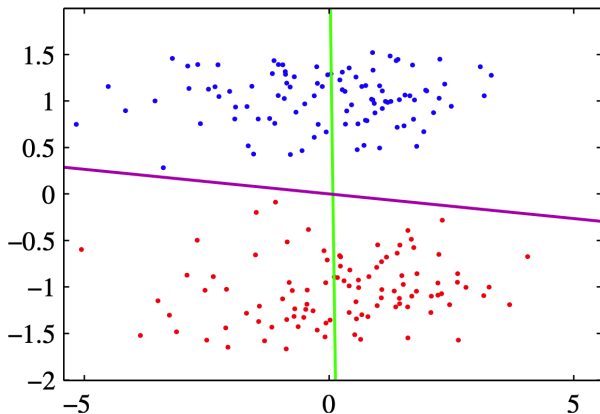


Figure: (Bishop Fig. 12.7) The first PC isn't always best for discrimination.

## Section 8

# TODO Summary

# Summary